

CSC413/2516 Winter 2023 — Course Information

Neural Networks and Deep Learning

Course web site: <http://uoft-csc413.github.io/2023>

Overview

It is very hard to hand design programs to solve many real world problems, e.g. distinguishing images of cats v.s. dogs. Machine learning algorithms allow computers to learn from example data, and produce a program that does the job. Neural networks are a class of machine learning algorithm originally inspired by the brain, but which have recently have seen a lot of success at practical applications. They're at the heart of production systems at companies like Google and Facebook for image processing, speech-to-text, and language understanding. This course gives an overview of both the foundational ideas and the recent advances in neural net algorithms.

Roughly the first 2/3 of the course focuses on supervised learning — training the network to produce a specified behavior when one has lots of labeled examples of that behavior. The last 1/3 focuses on unsupervised learning and reinforcement learning.

Schedule

This year, there are four sections: two afternoon sections and two night sections for the course. All sections will cover the same material, and will have the same assignments. See the course web page for detailed times. Tutorials are held after the lectures, which means there is no tutorial for the Lec0201/2101 section in Week 1.

Prerequisites

This is a second course in machine learning, so it has some substantial prerequisites. These prerequisites will be enforced, including for grad students.

- **Multivariable Calculus:**

MAT235/MAT237/MAT257/MAT291/MAT294/AER210/MAT232/MAT233/MATB41

- **Linear Algebra:**

MAT221H1/MAT223H1/MAT240H1/MAT185/MAT188/MAT223/MATA23

- **Machine Learning:**

CSC311/CSC411/STA314/ECE421/ROB313

Load

There are 24 hours of lectures and 11 hours of tutorials.

Readings

There is no required textbook for the class. A few small readings may be assigned if the need arises. These required readings will all be available on the web, for free.

There are also some relevant resources which are freely available online. We will try to provide links on a lecture-by-lecture basis.

- Video lectures for UofT Professor Geoffrey Hinton's Coursera course. Professor Hinton is one of the fathers of the field, so think of these as the Feynman Lectures of neural nets.
https://www.youtube.com/playlist?list=PLoRl3Ht4JOcdU872GhiYWf6jwrk_SNhz9
- *Deep Learning*, a textbook by Yoshua Bengio, Ian Goodfellow, and Aaron Courville.
<http://www.deeplearningbook.org/>
- Andrej Karpathy's lecture notes on convolutional networks. These are very readable and cover the material in roughly the first half of the course.
<http://cs231n.github.io/>
- Richard Socher's lecture notes, focusing on RNNs.
<http://cs224d.stanford.edu/syllabus.html>
- *Metacademy*, an online website (which one of the instructors is involved with) which helps you construct personalized learning plans and which has links to lots of resources relevant to particular concepts. We'll post links to relevant Metacademy concepts as the course progresses.
<http://www.metacademy.org>
- Video lectures for Hugo Larochelle's neural networks course. These are similar to Professor Hinton's lectures but a bit more mathematical.
http://info.usherbrooke.ca/hlarochelle/neural_networks/content.html
- David MacKay's excellent textbook, *Information Theory, Inference, and Learning Algorithms*. This isn't focused on neural nets per se, but it has some overlap with this course, especially the lectures on Bayesian models.
<http://www.inference.phy.cam.ac.uk/mackay/itila/>
- *Neural Networks and Deep Learning*, a book by physicist Michael Nielsen which covers the basics of neural nets and backpropagation.
<http://neuralnetworksanddeeplearning.com/>

Marking Scheme

The undergraduate marking scheme is as follows:

- Final project: 30%.
- Assignments: 60%
 - Total of 4, weighted equally.
 - **Important:** A minimum of 3 out of 4 assignments must be submitted on time (with grace days) to pass the course. See details below.
- Midterm test: 10%.

The marking scheme for graduate students is identical to the marking scheme for undergraduates, except that the midterm test exam is weighted towards the final project. The requirements and marking scheme for the final project will be posted separately on the course web page.

Academic Integrity

By the time you get to an advanced course like CSC413 you've heard this lots of times, so we'll keep it brief: avoid academic offenses (a.k.a. cheating). All graded work in this course is individual work except for the final project.

Assignments

Each assignment is consist of two parts: written and programming exercises.

- **Written questions**

In order to give you additional practice with the material, we assign written homeworks, which give you additional practice with the course content and encourage you to keep on top of the material. Roughly speaking, there will be one homework due each week that doesn't have another assignment or test. Each one consists of 2-3 conceptual questions and is meant to take a few hours.

- **Programming questions**

A typical assignment will require you to write (or modify) and use some Python code that implements a simple version of a learning procedure that has recently been covered in the course. You will have to submit a brief report (roughly two pages plus figures) that describes the results you obtained.

Dates. Assignments will typically be due at 11:59pm on Friday. See the course web page for particular deadlines. Each homework covers material up through the lecture one week prior to the deadline.

Format. Assignments must be submitted in PDF and the corresponding iPython Notebooks format through MarkUs. We encourage typesetting using L^AT_EX, but scans of handwritten solutions are also acceptable.

Lateness and grace days. Every student has a total of 7 grace days to extend the coursework deadlines through the semester. Each grace day allows for a 24 hours deadline extension without late penalty. That is, you may apply the grace days on a late submission to remove its late penalty. The maximum extension you may request is up to the remaining grace days you have banked up. We will keep track of the grace days on MarkUs. After the grace period, assignments will be accepted up to 3 days late, but 10% will be deducted for each day late, rounded up to the nearest day. After that, submissions will not be accepted and will receive a score of 0, and thus deemed NOT “on time.” There is no other extension mechanism outside of the grace days.

An assignment is considered “on time” when the submission is accepted on MarkUs within the application of grace days.

Weighting. In aggregate, the assignments count for 60% of the total grade for the course, so individually they count for roughly 15% each.

Collaboration policy. You are expected to work on the assignments by yourself. You should not discuss them with anyone except the tutors or the instructor. The report you hand in should be entirely your own work and you may be asked to demonstrate how you got any results that you report.

Computation Resources

Many of the deep learning success stories in the recent years rely on the advances of modern GPU computing. The programming assignments here are lightweight comparing to the state of the art deep learning models in terms of their computation requirement. But we highly recommend you to debug your models and to complete the experiments on a modern GPU. Here are the list of free computation resources you have access to:

Colab (Recommended) Google Colab is a web-based iPython Notebook service that has access to a free Nvidia K80 GPU per Google account. Although it was initially developed for TensorFlow usage, Colab can easily be configured to run PyTorch, see tutorial here: https://pytorch.org/tutorials/beginner/basics/quickstart_tutorial.html and click “Run in Google Colab” on the top of the page.

GCE (Recommended) Google Compute Engine delivers virtual machines running in Google’s data center. You get \$300 free credit when you sign up. They provide some of the latest GPUs on the market.

AWS-EC2 Amazon Elastic Compute Cloud (EC2) is a popular cloud platform. You may get free credit somewhere online.

CS Teaching Lab There are some very old GPUs in our CS Teaching Labs / CDF labs, see <https://www.teach.cs.toronto.edu/faq.html#ABOUT5> for details.

Tests

Midterm. The midterm test (worth 10% of the course grade) will be held on Thursday Feb 9th. It covers material up through Lecture 4 (one week prior to the test). The midterm test will be accessible on Quercus for a 24-hour duration.

Missed tests. Missed test will get a score of 0. In the event of illness, students should contact us at ticket-csc413-2023-01@teach.cs.toronto.edu before the test date and approved by the instructor. We will arrange a make-up test within 7 days of the original midterm time.

Online Forum

We'll use Piazza for the course forum, <http://piazza.com/utoronto.ca/winter2023/csc4132516>. We expect Piazza to be the primary source of your interactive learning experience.

Accessibility

The CSC413 teaching staff is fully committed to ensuring accessibility for all our students. For the students looking for additional academic accommodations or accessibility services registration, please visit www.accessibility.utoronto.ca. Students are encouraged to review the course syllabus at the beginning of a course and discuss questions regarding their accommodations for the course with their Accessibility Advisor. Once registered, students should send the Letter of Academic Accommodations to our ticketing system at ticket-csc413-2023-01@teach.cs.toronto.edu as soon as possible by Friday, January 27th, 2023.

Waitlist and Course Enrollment

CSC413 and CSC2516 always had long waiting lists for the last few years. The hard enrollment cap is determined by teaching resources available at the department level. Note that waitlists typically expire one week after the course starts. Once waitlists are removed, students are responsible for trying to enroll in the course on ACORN in a first-come, first-serve fashion. If you have further questions, please get in touch with [CS undergrad office](#) or [CS graduate office](#).

Auditing

If you are not registered in the class, it is possible for you to audit it (sit in on the lectures). Here are the official university rules on auditors (taken from the Department of Computer Science instructor's advice page):

To audit a course is to sit and listen to the lectures, and perhaps to the tutorials, without formally enrolling. Auditing is acceptable *if the auditor is a student at U of T, and no University resources are to be committed to the auditor*. The “must be a student” condition means that students of other universities, employees of outside organizations (or even of U of T itself!), or any other non-students, are not permitted to be auditors. (If we did not have this rule, the University would require us to collect auditing fees, and we are not willing to do that.)

The “no resources used” condition means that auditors do not get computing accounts, cannot have term work marked, and cannot write exams. In other words, they cannot use instructors time, TA time, or administrative resources of any kind.

An auditor may not attend class *unless there is an empty seat after the last regularly-enrolled student has sat down*. That sounds frivolous, but in fact it is an aspect of an important point: if enrollment in a course has been closed because the room size has been reached, then there may well be physical seats for auditors, because it is rare for every student to appear for a lecture, but auditors will not be allowed to enroll later on in the course, even if some students drop it. Neither instructors nor the department can waive this rule.

Often these conditions are perfectly acceptable to auditors; we don’t mean to ban the practice, but only to live within the University’s rules.